

STA108 Term Project

By Hallie Pimperl and Rohit Raghuraman Rajagopalan

Appendix

Term Project





```
#Subset 80% of data
countries_data <- read.csv("/Users/hpimp_girlpower/Downloads/countries.csv")
n_count= nrow(countries_data)
set.seed(10)
subset_id_count= sample(n_count,0.8*n_count)
count_subset= countries_data[subset_id_count, ]

#Set relevant variables
area=count_subset$LandArea
pop=count_subset$Population
rural=count_subset$Rural
health=count_subset$Health
internet=count_subset$Internet
elderly=count_subset$ElderlyPop
birth=count_subset$BirthRate
life_expect=count_subset$LifeExpectancy
carbon=count_subset$CO2
GDP=count_subset$GDP
cell=count_subset$Cell
```





```
#Load needed libraries
library(MASS)
library(leaps)
```

Initial Relationship



```
#Scatterplots for each predictor to determine relationship to outcome variable
par(mfrow=c(2,2))
plot(area, life_expect,xlab= 'Land Area',ylab= 'Life Expectancy', main='Area vs. Life Expectancy')
plot(pop,life_expect,xlab= 'Population',ylab= 'Life Expectancy', main='Population vs. Life Expectancy')
plot(rural,life_expect,xlab= 'Rural Percentage',ylab= 'Life Expectancy', main='Rural vs. Life Expectancy')
plot(health,life_expect,xlab= 'Health',ylab= 'Life Expectancy', main='Health vs. Life Expectancy')
```



```
plot(internet,life_expect,xlab= 'Internet Access',ylab= 'Life Expectancy', main='Internet vs. Life
Expectancy')
plot(elderly,life_expect,xlab= 'Elderly Population',ylab= 'Life Expectancy', main='Elderly vs. Life Expectancy')
plot(birth,life_expect,xlab= 'Birth Rate',ylab= 'Life Expectancy', main='Birth vs. Life
Expectancy')
plot(carbon,life_expect,xlab= 'Carbon Emissions',ylab= 'Life Expectancy', main='Carbon vs. Life
Expectancy')
```



```
plot(GDP,life_expect,xlab= 'GDP',ylab= 'Life Expectancy', main='GDP vs. Life
Expectancy')
plot(cell,life_expect,xlab= 'Cell Phone',ylab= 'Life Expectancy', main='Cell vs. Life Expectancy')
```



Simple Regression Diagnostics

#Run simple linear regressions for each predictor to check for assumption violations





```
par(mfrow=c(2,2))
```

#AREA

```
area_model= lm(life_expect~area)
plot(area_model,which=1,main='Area')
plot(area_model,which=2,main='Area')
```





#POPULATION

```
pop_model= lm(life_expect~pop)
plot(pop_model,which=1,main='Population')
plot(pop_model,which=2,main='Population')
```







```
#RURAL
rural_model= lm(life_expect~rural)
plot(rural_model,which=1,main='Rural')
plot(rural_model,which=2,main='Rural')

#HEALTH
health_model= lm(life_expect~health)
plot(health_model,which=1,main='Health')
plot(health_model,which=2,main='Health')
```







```
#INTERNET
internet_model= lm(life_expect~internet)
plot(internet_model,which=1,main='Internet')
plot(internet_model,which=2,main='Internet')

#ELDERLY
elderly_model= lm(life_expect~elderly)
plot(elderly_model,which=1,main='Elderly')
plot(elderly_model,which=2,main='Elderly')
```




```
#BIRTH
birth_model= lm(life_expect~birth)
plot(birth_model,which=1,main='Birth')
plot(birth_model,which=2,main='Birth')

#CARBON
carbon_model= lm(life_expect~carbon)
plot(carbon_model,which=1,main='Carbon')
plot(carbon_model,which=2,main='Carbon')
```



```
#GDP
GDP_model= lm(life_expect~GDP)
plot(GDP_model,which=1,main='GDP')
plot(GDP_model,which=2,main='GDP')

#CELL
cell_model= lm(life_expect~cell)
plot(cell_model,which=1,main='Cell')
plot(cell_model,which=2,main='Cell')
```

Almost all plots show a deviation from normality and/or nonconstant errors, indicating a transformation in Y is needed. The variables: area, pop, internet, elderly, carbon, GDP, and cell all show signs of nonlinearity, indicating a transformation in X is also needed.

Initial Multiple Linear Regression

#First look at which relationships are significant using t-test

```
overall_model= lm(life_expect~area+pop+rural+health+internet+elderly+birth+carbon+GDP+cell)
summary(overall_model)
```

Call:

```
lm(formula = life_expect ~ area + pop + rural + health + internet + elderly + birth + carbon + GDP + cell)
```

Residuals:

```
Min 1Q Median 3Q Max
-16.4565 -2.3102 0.1329 3.1021 11.0534
```

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	8.108e+01	3.470e+00	23.362	< 2e-16 ***
area	-2.933e-07	2.399e-07	-1.223	0.22352
pop	1.008e-03	4.386e-03	0.230	0.81853
rural	-2.069e-02	2.633e-02	-0.786	0.43334
health	1.202e-01	1.119e-01	1.074	0.28459
internet	9.961e-02	3.560e-02	2.798	0.00588 **


elderly	-4.164e-01	1.632e-01	-2.551	0.01184 *
birth	-6.857e-01	7.749e-02	-8.849	3.96e-15 ***
carbon	-1.256e-01	8.147e-02	-1.541	0.12559
GDP	5.059e-05	4.974e-05	1.017	0.31089
cell	2.826e-02	1.354e-02	2.086	0.03879 *

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
 Residual standard error: 4.94 on 137 degrees of freedom
 Multiple R-squared: 0.7881, Adjusted R-squared: 0.7727
 F-statistic: 50.96 on 10 and 137 DF, p-value: < 2.2e-16

According to the above p-values, internet, elderly, birth, and cell are significant predictors of life expectancy.

#More regression diagnostics to see if violations were corrected

```
par(mfrow=c(1,2))
plot(overall_model,which=1)
plot(overall_model,which=2)
```



Indications of nonlinearity, nonconstant variance, and non-normality still exist. Transform variables.

Y Transformations

#Use Box-Cox Method to evaluate best y-transformation

```
par(mfrow=c(2,5))
```

```
#bc_area = boxcox(area_model, lambda=seq(-10,10,1/10), plotit=F)
```

```

bc_pop=   boxcox(pop_model,lambda= seq(-7,7,1/10),plotit=F) bc_rural=
          boxcox(rural_model,lambda= seq(-7,7,1/10),plotit=F)
bc_internet=   boxcox(internet_model,lambda= seq(-7,7,1/10),plotit=F) bc_health=
               boxcox(health_model,lambda= seq(-7,7,1/10),plotit=F) bc_elderly=
               boxcox(elderly_model,lambda= seq(-7,7,1/10),plotit=F) bc_birth=
               boxcox(birth_model,lambda= seq(-7,7,1/10),plotit=F) bc_carbon=
               boxcox(carbon_model,lambda= seq(-7,7,1/10),plotit=F) bc_GDP=
               boxcox(GDP_model,lambda= seq(-7,7,1/10),plotit=F)
bc_cell=      boxcox(cell_model,lambda= seq(-7,7,1/10),plotit=F)

#bc_area$[which.max(bc_area$y)]

cbind(bc_pop$x[which.max(bc_pop$y)], bc_rural$x[which.max(bc_rural$y)],
      bc_internet$x[which.max(bc_internet$y)], bc_health$x[which.max(bc_health$y)],
      bc_elderly$x[which.max(bc_elderly$y)])

```

```

[,1] [,2] [,3] [,4] [,5]
[1,]3.2   3.9   5.3   3.5   4

```

As we can see, there is a common pattern of the range of lambda values that fall within the 95% confidence interval. Therefore, I will choose a lambda=4 since it is a value that occurs often.

X Transformations

Transformations were found through trial-and-error. Examples of their transformed relationships with our transformed life expectancy are shown below.


```

par(mfrow=c(1,2))


plot(log(area), life_expect**4,xlab= 'Log(Area)',ylab= 'Life Expectancy**4', main='Area v. Life Expectancy')
plot(log(elderly), life_expect**4,xlab= 'Log(Elderly)',
      ylab='Life Expectancy**4',main= 'Elderly v. Life Expectancy')

```

Area v. Life Expectancy




Elderly v. Life Expectancy




```
plot(log(GDP), life_expect**4,xlab= 'Log(GDP)',ylab= 'Life Expectancy**4', main='GDP v. Life
Expectancy')
plot(log(cell), life_expect**4,xlab= 'Log(Cell)',ylab= 'Life Expectancy**4', main='Cell v. Life Expectancy')
```

GDP v. Life Expectancy



Cell v. Life Expectancy



Transformed Multiple Linear Regression

```
new_LE=life_expect    **4
new_area= log(area)
new_pop= log(pop)
new_internet=internet  **(1/2)
new_elderly= log(elderly)
new_GDP= log(GDP)
new_cell= log(cell)
```

#Throw transformed variables into new MLR to see which predictors are significant now

```
transformed_model= lm(new_LE~new_area+new_pop+rural+new_internet+health+
                      new_elderly+birth+new_GDP+new_cell)
summary(transformed_model)
```

Call:

```
lm(formula = new_LE ~ new_area + new_pop + rural + new_internet + health +
    new_elderly + birth + new_GDP + new_cell)
```

Residuals:

```
Min    1Q  Median    3Q   Max
-15396594 -2786981 -203958  3407028 13838525
```

Coefficients:

```
Estimate Std. Error t value Pr(>|t|)
```

(Intercept)	23010024	7835143	2.937	0.00389	**
new_area	-879946	287339	-3.062	0.00264	**
new_pop	820087	369060	2.222	0.02790	*
rural	-35053	28364	-1.236	0.21862	
new_internet	1474187	356962	4.130	6.26e-05	***
health	160069	114099	1.403	0.16289	
new_elderly	-42689	1135620	-0.038	0.97007	
birth	-409824	87608	-4.678	6.83e-06	***
new_GDP	1417013	648873	2.184	0.03067	*
new_cell	93080	659389	0.141	0.88795	

 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
 Residual standard error: 5018000 on 138 degrees of freedom
 Multiple R-squared: 0.8325, Adjusted R-squared: 0.8215
 F-statistic: 76.19 on 9 and 138 DF, p-value: < 2.2e-16

Now, according to our p-values, area, population, internet, birth, and GDP are significant predictors of life expectancy.

Final Model

#Attempt model with only important variables

```
final_model= lm(life_expect~4~new_area+new_pop+new_internet+birth+new_GDP)
summary(final_model)
```

Call:

```
lm(formula = life_expect^4 ~ new_area + new_pop + new_internet + ##      birth +
new_GDP)
```

Residuals:


Min	1Q	Median	3Q	Max
-15089880	-3064764	-134918	3297799	13817901

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	18420514	5403623	3.409	0.000849	***
new_area	-826711	284426	-2.907	0.004241	**
new_pop	751206	360547	2.084	0.038997	*
new_internet	1675575	322733	5.192	7.07e-07	***
birth	-390819	67880	-5.757	5.07e-08	***
new_GDP	1804916	560596	3.220	0.001591	**

 Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
 Residual standard error: 5008000 on 142 degrees of freedom
 Multiple R-squared: 0.8283, Adjusted R-squared: 0.8222
 F-statistic: 137 on 5 and 142 DF, p-value: < 2.2e-16

```
par(mfrow=c(1,2))
plot(final_model,which=1)
plot(final_model,which=2)
```



Residual and QQplot show no strong violations of our assumptions. All of our predictors are significant.

Model Selection

#Use forward/backward selection to understand most important predictors

```
forward_LEmodel= regsubsets(new_LE~new_area+new_pop+new_internet+birth+
                           new_GDP,data=count_subset,method="forward")
summary(forward_LEmodel)
```

Subset selection object

Call: `regsubsets.formula(new_LE ~ new_area + new_pop + new_internet + ## birth + new_GDP, data = count_subset, method = "forward")`

5 Variables (and intercept)

Forced in Forced out

new_area FALSE FALSE

new_pop FALSE FALSE

new_internet FALSE FALSE

birth FALSE FALSE

new_GDP FALSE FALSE

1 subsets of each size up to 5 ## Selection

Algorithm: forward

	new_area	new_pop	new_internet	birth	new_GDP
1	(1) " " " "	"*"	" "	" "	
2	(1) " "	" "	"*"	"*"	" "
3	(1) " "	" "	"*"	"*"	"*"
4	(1) "*"	" "	"*"	"*"	"*"
5	(1) "*"	"*"	"*"	"*"	"*"

Backward model selection shows the same results.

According to forward selection, the most significant variables in order from most significant to least significant are: internet, birth, GDP, area, and population.

#Use LEAPS to do model selection to pick a parsimonious model

```
library(leaps)
design_final=      cbind(new_area, new_pop, new_internet, birth, new_GDP)
```

#try it with adjusted R2

```
adjr2_model=      leaps(x=design_final,y=new_LE,method=      c('adjr2'),
                        names=c('Area','Population','Internet','Birth','GDP'))
adjr2_model$adjr2
```

```
[1] 0.726358152 0.702230338 0.684551562 0.070835858 0.004828808 0.804816576
[7] 0.777720664 0.765348990 0.738928449 0.727276815 0.705900683 0.700491715
[13] 0.693500283 0.682776915 0.093910913 0.814170905 0.809316803 0.804212790
[19] 0.780584413 0.776364362 0.773879336 0.763967551 0.741346056 0.716607271
[25] 0.708139311 0.818089341 0.812983820 0.810599851 0.789978617 0.782279844
[31] 0.822242449
```

The highest parsimonious adjusted R2 is with population excluded.

#Try with Mallow's



```
mallow_LEmodel=      leaps(x=design_final,y=new_LE,method=      c('Cp'),
                           names=c('Area','Population','Internet','Birth','GDP'))
mallow_LEmodel$Cp
```

```
[1] 80.753939 100.571161      115.091508 619.162881 673.377340 17.214595
[7] 39.317213      49.409008      70.960713 80.465159 97.902051 102.314242
[13] 108.017278 116.764520      597.113005 10.538695 14.470965 18.605686
[19] 37.746849      41.165479      43.178579 51.208039 69.533535 89.574230
[25] 96.434059      8.341037      12.448257 14.366080 30.955173 37.148578
[31] 6.000000
```

Best parsimonious value of Mallow's Cp is again with population excluded.

#final model results





```
final_model=      lm(new_LE~new_internet+birth+new_GDP+new_area)
par(mfrow=c(1,2))
plot(final_model,which=1) plot(final_model,which=2)
```

```
standard_model= lm(scale(new_LE)~scale(new_internet)+scale(birth)+
                  scale(new_GDP)+scale(new_area))
confint(standard_model)
```

	2.5 %	97.5 %
(Intercept)	-0.06930064	0.069300643
scale(new_internet)	0.24977023	0.526902541
scale(birth)	-0.48325676	-0.243737786
scale(new_GDP)	0.05989644	0.341165992
scale(new_area)	-0.14532020	-0.001764897


```
par(mfrow=c(2,2))
plot(new_internet, new_LE,xlab= 'Internet Access',ylab= 'Life Expectancy', main='Internet
Access')
plot(birth, new_LE,xlab= 'Birth Rate',ylab= 'Life Expectancy', main='Birth Rate')
plot(new_GDP, new_LE,xlab= 'GDP',ylab= 'Life Expectancy', main='GDP')
plot(new_area, new_LE,xlab= 'Land Area',ylab= 'Life Expectancy', main='Land Area')
```



Multicollinearity

#Check for multicollinearity for all variables

```
pairs(cbind(area,pop,rural,health,internet,elderly,birth,carbon,GDP,cell))
```



#Displaying correlation matrix of all the predictors

```
design = cbind(life, area, pop, elderly, cell, rural, health, birth, internet, carbon, gdp)
cor(design)
```

(Below result is the edited version of the matrix to accommodate all the values)

	life	area	popu	eldpopu	cell	rural	health	birth	internet	carbon	gdp
life	1	0.00718	-0.0249	0.641	0.685	-0.621	0.265	-0.850	0.7414	0.431	0.617
area		1	0.342	0.0885	0.047	-0.131	0.00749	-0.067	0.741	0.431	0.616
popu			1	-0.013	-0.0639	0.0636	-0.1445	-0.0115	0.077	0.128	0.077
elderly				1	0.495	-0.495	0.393	-0.75	-0.039	-0.034	-0.036
cell					1	-0.60	0.197	-0.684	0.745	0.205	0.56
rural						1	-0.191	0.599	0.597	0.443	0.49
health							1	-0.1996	-0.665	-0.43	-0.649
birth								1	0.395	0.027	0.343
internet									1	0.4855	0.804
carbon										1	0.592
gdp											1

#Displaying Variance Inflation Factor (VIF) of all the predictors


```
library(faraway)
vifmodel <- lm(life ~ area + health + carbon + elderly + popu + rural + cell + birth + gdp + internet)
vif(vifmodel)
```

```
area health carbon elderly popu rural cell birth gdp internet
1.203130 1.365423 2.063836 3.788557 1.197367 2.302427 2.208491 4.037504 3.736785 5.051648
```

Internet, GDP, Birth rate are moderately collinear but Area is less collinear with other predictors.

#Check for multicollinearity for selected predictors

```
pairs(cbind(internet, birth, GDP, area))
```



Multicollinearity exists between internet and birth, birth and GDP, and internet and GDP.

#Displaying correlation matrix of the selected predictors

```
Design1 = cbind(internet,birth,gdp,area)
cor(design1)
```

	internet	birth	gdp	area
internet	1.00000000	-0.71589970	0.80441383	0.07777323
birth	-0.71589970	1.00000000	-0.57035251	-0.06777425
gdp	0.80441383	-0.57035251	1.00000000	0.07770053
area	0.07777323	-0.06777425	0.07770053	1.00000000

#Displaying Variance Inflation Factor (VIF) of the model predictors

```
vifmodel1 <- lm(life ~ area + birth + gdp + internet)
vif(vifmodel1)
```

	area	birth	gdp	internet
	1.007060	2.052343	2.835907	3.922992

Before transformations, Internet has the highest VIF value of 3.92, followed by GDP, birth rate and area. Area is the least correlated and Internet is moderately correlated with other model predictors.

#Displaying Variance Inflation Factor (VIF) of the transformed model predictors

```
vifmodel1 <- lm(life^4 ~ log(area) + birth + log(gdp) + sqrt(internet))
vif(vifmodel1)
```

	log(area)	birth	log(gdp)	sqrt(internet)
	1.065515	2.966203	4.090407	3.970959

After transformation, GDP has the highest VIF value of 4.09, which is not so different from that of Internet (3.97), and followed by birth rate and area. Hence, GDP and Internet are moderately correlated whereas birth rate is less correlated, and area is the least correlated.